



Hochschule für Angewandte Wissenschaften Hamburg  
*Hamburg University of Applied Sciences*

## Optimizer / Statistiken / SQL / Performance

**Marcus Bender**

STU (Strategisch Technische Unterstützung)

Oracle Deutschland

[marcus.bender@oracle.com](mailto:marcus.bender@oracle.com)



# Überblick Performance

- 1 HTTP Request / Netzwerk
- 2 Applikation Server / Java Klassen / Netzwerkanbindung
- 3 Optimizer / SQL Tuning
- 4 Datenmodellierung / Partitionierung / Indizes
- 5 Datenbanktuning
- 6 System / OS / VMware / Netzwerk / IO-Subsystem



# Applikation / Charakteristik der Anwendung

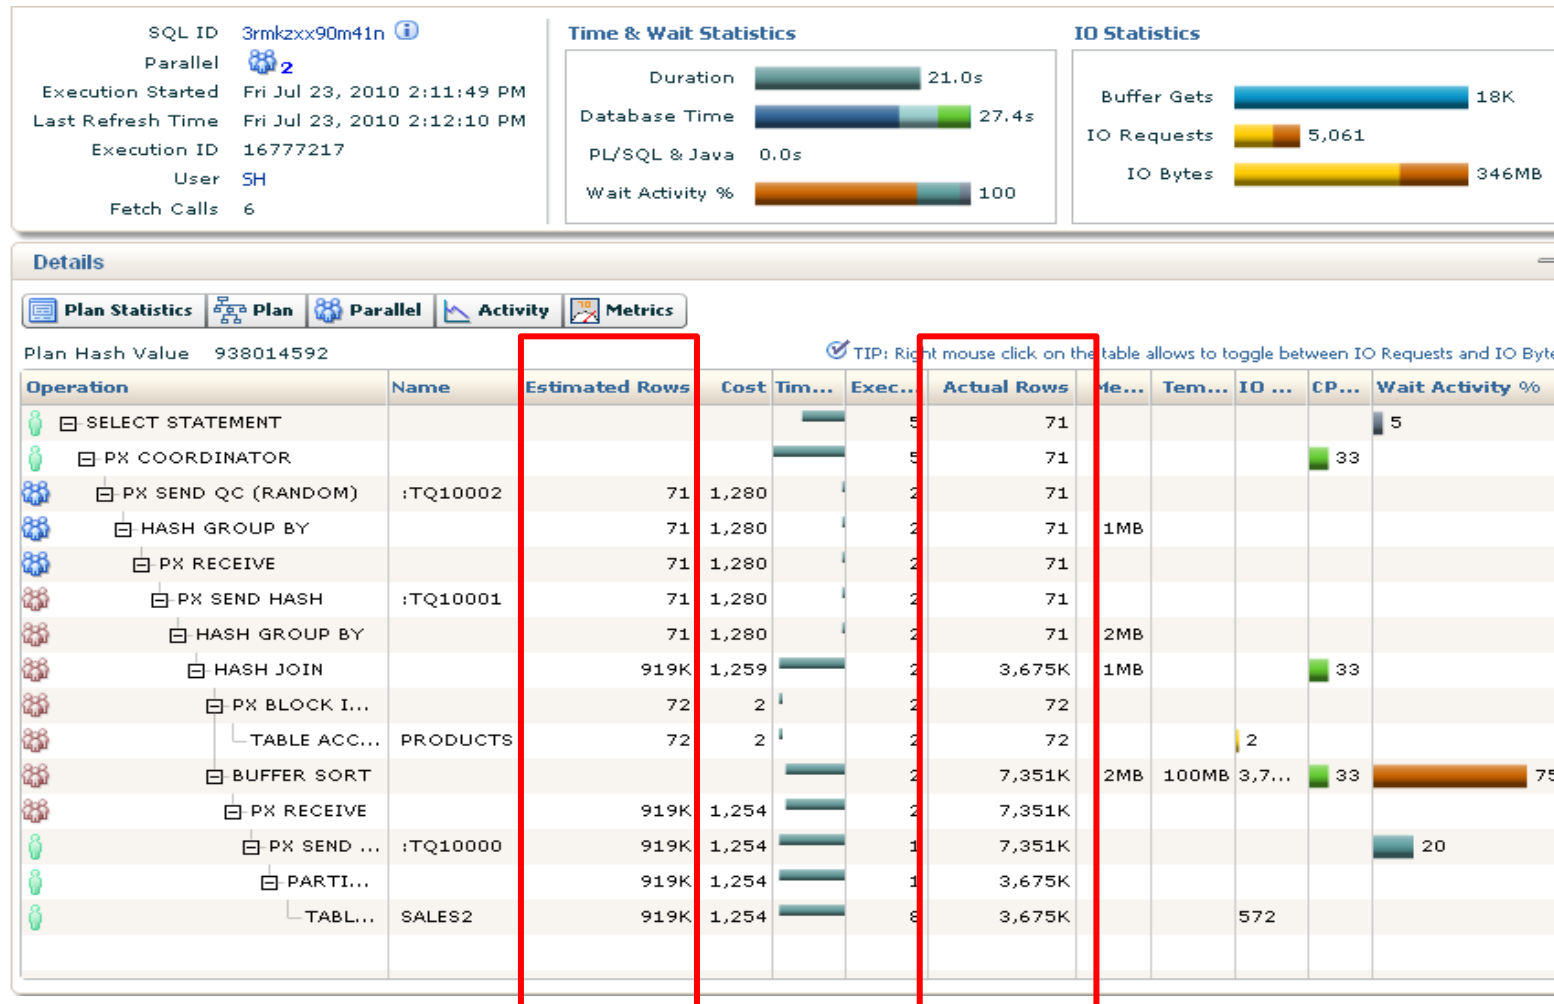
OLTP / Einzelsatzverarbeitung	Data Hub / ODS / Mischbetrieb	DWH / Massendatenverarbeitung
<ul style="list-style-type: none"><li>• Viele Anwender</li><li>• Tausende Transaktionen pro Sekunde, sehr viele Ausführungen</li><li>• Einzelnes Statement deutlich weniger als eine Sekunde</li><li>• Einzelsatzverarbeitung Row-by-row based Cursorverarbeitung Nested Loop</li><li>• Pro Ausführung werden kleine Datenmengen verarbeitet</li><li>• B*-Tree Indizes! Single Read/Write I/O!</li><li>• Partitionierung? / Parallel Query?</li></ul>	<ul style="list-style-type: none"><li>• Viele Anwender bzw. Datenabnehmer</li><li>• ETL Jobs, Reporting, OLTP, viele &amp; anspruchsvolle Ausführungen</li><li>• Unterschiedliche lange Ausführungszeiten</li><li>• Einzel- + Massenverarbeitung Row- + Set based Alle Aufrufarten NL + Hash-Join</li><li>• Pro Ausführung werden unterschiedl. große Datenmengen verarbeitet</li><li>• Indizes! Single Read/Write I/O!</li><li>• Partitionierung! / Parallel Query! IO-Durchsatz! Star-Datenmodell! Bitmap Indizes</li></ul>	<ul style="list-style-type: none"><li>• Wenige Anwender</li><li>• ETL Jobs, Reporting, wenige anspruchsvolle Ausführungen</li><li>• Einzelnes Statement kann Minuten, bzw. Stunden dauern</li><li>• Massenverarbeitung Set based Direkter SQL Aufruf Merge-Join / Hash-Join</li><li>• Pro Ausführung werden große Datenmengen verarbeitet</li><li>• Indizes? Single Read/Write I/O?</li><li>• Partitionierung! / Parallel Query! Full Table Scan! IO-Durchsatz! Star-Datenmodell! Bitmap Indizes</li></ul>



# SQL Statements / Monitoring im EM Cloud Control

## SQL Monitor

- Gibt einen exakten Überblick über laufende SQL Statements
- Execution Plan
- Anzahl gelesener Sätze
- Anzahl I/Os
- Wait Aktivitäten
- Parallele Verarbeitung
- Auch für Entwickler wichtig!

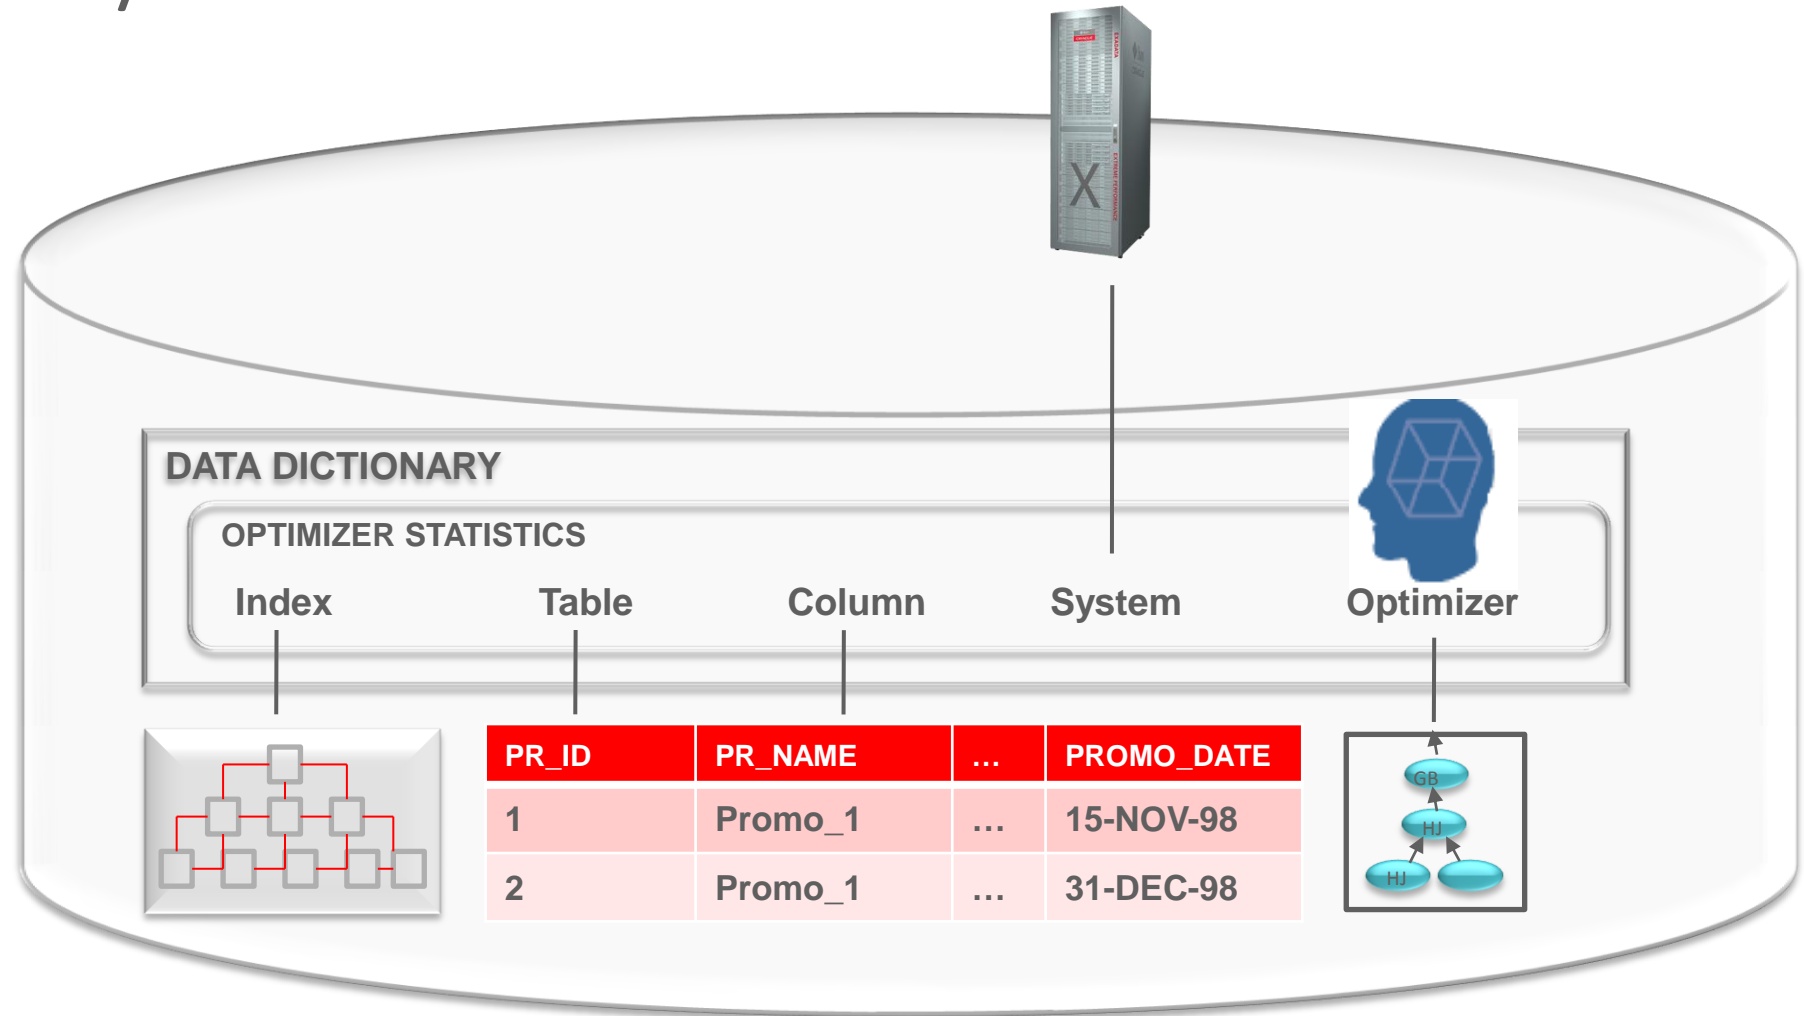


# SQL Statements / Statistiken



## Statistiken

- Tabellenstatistik
  - Größe
  - Anzahl Sätze
  - Satzlänge
- Spaltenstatistik
  - Kardinalität
  - Verteilung
  - Häufigkeit
- Indexstatistik
  - Indextiefe
  - Indexebenen
- Systemstatistiken
  - Anzahl CPUs
  - IO Durchsatz
  - Single Block Read

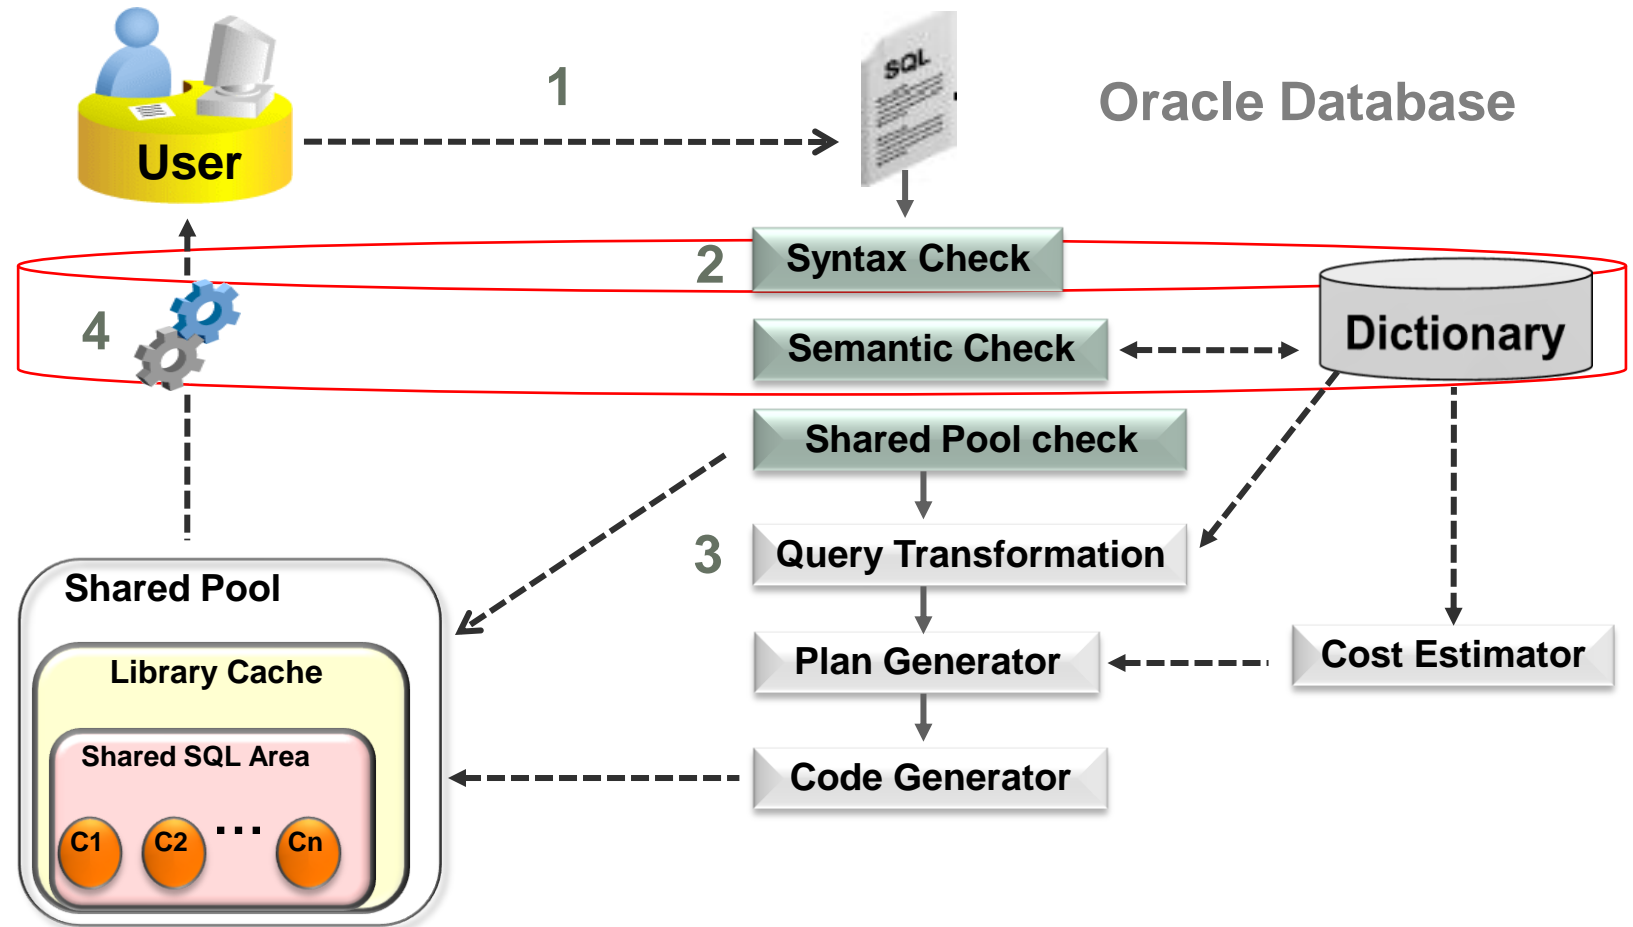




# SQL Statements / Arbeitsweise des Optimizers

## Optimizer / Execution Plan

1. Ein User setzt ein SQL Statement ab
2. Soft Parse: des Statements mit Hilfe des Dictionaries und des Shared Pools
3. Hard Parse: es wird mit Hilfe von Statistiken der beste Execution Plan gesucht und als ausführbarer Code in den Shared Pool geschrieben
4. Ausführung des Statements und senden des Ergebnisses an den User






# SQL Statements / Arbeitsweise des Optimizers

## Optimizer / Execution Plan

- Kardinalität
  - Einschätzung der Selektivität / Kardinalität
  - Hohe Selektivität (Auftragsnummer), Geringe Selektivität (Geschlecht)
- Zugriffspfade
  - Bestimmung des Zugriffs auf die Daten
  - Full Table Scan, Index, Bitmap Index, Index Range Scan, ...
- Join-Typ
  - Welche Join Art wurde verwendet
  - Nested Loop Join, Hash Join, Sort-Merge Join, Kartesisches Produkt
- Join-Reihenfolge
  - Bestimmung der Reihenfolge
  - Erst A mit B joinen, dann mit D, zuletzt mit C





# SQL Statements / Arbeitsweise des Optimizers

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				709 (100)			
1	VIEW		4	156	709 (34)	00:00:09		
2	SORT GROUP BY		4	380	709 (34)	00:00:09		
* 3	HASH JOIN		4	380	708 (34)	00:00:09		
* 4	HASH JOIN <b>Join-Typ</b>		4	276	701 (34)	00:00:09		
5	PART JOIN FILTER CREATE	:BF0000	4	156	35 (18)	00:00:01		
6	VIEW	VW_GBF_10	4	156	35 (18)	00:00:01		
7	HASH GROUP BY		4	208	35 (18)	00:00:01		
8	PARTITION RANGE SINGLE		4	208	34 (15)	00:00:01	14	14
9	PARTITION HASH ALL		4	208	34 (15)	00:00:01	1	128
* 10	<b>1</b> TABLE ACCESS FULL	TRANS_DETAIL	4	208	34 (15)	00:00:01	1665	1792
11	PARTITION RANGE ALL		130K	3829K	655 (34)	00:00:08	1	16
12	PARTITION HASH JOIN-FILTER		130K	3829K	655 (34)	00:00:08	:BF0000	:BF0000
13	<b>2</b> TABLE ACCESS FULL <b>Pfad</b>	ACCOUNT_MASTER	130K	3829K	655 (34)	00:00:08	1	2048
14	<b>3</b> TABLE ACCESS FULL	PCODE_REF	1307	33982	6 (0)	00:00:01		

Predicate Information (identified by operation id):

**Kardinalität**

```
3 - access("A"."PCODE"="B"."PCODE")
4 - access("A"."ACCT_NUM"="ITEM_2" AND "A"."CO_ID"="ITEM_1")
10 - filter(("C"."ASOF_YYYYMM"=200102 AND "C"."TRAN_AMT"<2000000000))
```





# SQL Statements / Access Paths

Access Path	Explanation
Full table scan	Reads all rows from table & filters out those that do not meet the where clause predicates. Used when no index, DOP set etc
Table access by Rowid	Rowid specifies the datafile & data block containing the row and the location of the row in that block. Used if rowid supplied by index or in where clause
Index unique scan	Only one row will be returned. Used when stmt contains a UNIQUE or a PRIMARY KEY constraint that guarantees that only a single row is accessed.
Index range scan	Accesses adjacent index entries returns ROWID values Used with equality on non-unique indexes or range predicate on unique index (<.>, between etc)
Index skip scan	Skips the leading edge of the index & uses the rest Advantageous if there are few distinct values in the leading column and many distinct values in the non-leading column
Full index scan	Processes all leaf blocks of an index, but only enough branch blocks to find 1 <sup>st</sup> leaf block. Used when all necessary columns are in index & order by clause matches index struct or if sort merge join is done
Fast full index scan	Scans all blocks in index used to replace a FTS when all necessary columns are in the index. Using multi-block IO & can go parallel
Index joins	Hash join of several indexes that together contain all the table columns that are referenced in the query. Won't eliminate a sort operation
Bitmap indexes	uses a bitmap for key values and a mapping function that converts each bit position to a rowid. Can efficiently merge indexes that correspond to several conditions in a WHERE clause




# SQL Statements / Join Types

Access Path	Explanation
Nested Loops joins	For every row in the outer table, Oracle accesses all the rows in the inner table Useful when joining small subsets of data and there is an efficient way to access the second table (index look up)
Hash Joins	The smaller of two tables is scanned and resulting rows are used to build a hash table on the join key in memory. The larger table is then scanned, join column of the resulting rows are hashed and the values used to probe the hash table to find the matching rows. Useful for larger tables & if equality predicate
Sort Merge joins	Consists of two steps: 1. Sort join operation: Both the inputs are sorted on the join key. 2. Merge join operation: The sorted lists are merged together. Useful when the join condition between two tables is an inequality condition
Cartesian Joins	Joins every row from one data source with every row from the other data source, creating the Cartesian Product of the two sets. Only good if tables are very small. Only choice if there is no join condition specified in query
Outer Joins	Returns all rows that satisfy the join condition and also returns all of the rows from the table without the (+) for which no rows from the other table satisfy the join condition

# SQL Statements / Hash Join

- Für Batch Jobs / Massenverarbeitung
- Sinnvoll -> DWH
- Historie
  - Sort/Merge Joins (SMJ) und/oder
  - Nested Loop Joins nicht immer optimal
- Hash Join bietet ein Verfahren, das auch
- dann sehr effizient arbeitet, wenn keine
- Indizes vorhanden sind!



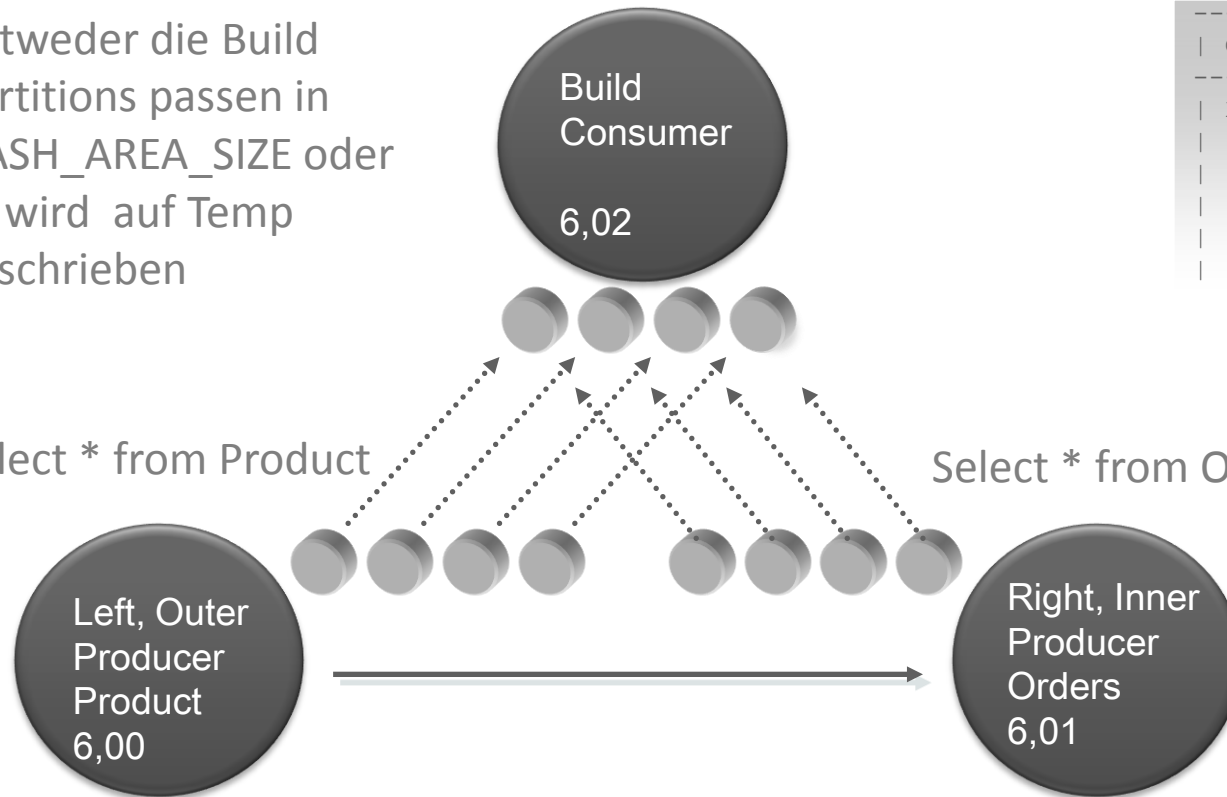


# SQL Statements / Parallel Hash Join

Entweder die Build  
Partitions passen in  
HASH\_AREA\_SIZE oder  
es wird auf Temp  
geschrieben

Select \* from Product

Select \* from Orders




Operation	Name	TQ	IN-OUT	PQ Dist
HASH JOIN		6,02	PCWP	
PARTITION HASH ALL		6,02	PCWP	
FULL TABLE SCAN	PRODUCT	6,00	P->P	HASH
PARTITION RANGE ALL		6,02	PCWP	
PARTITION HASH ALL		6,02	PCWP	
FULL TABLE SCAN	ORDERS	6,01	P->P	HASH



# SQL Statements / Cardinality

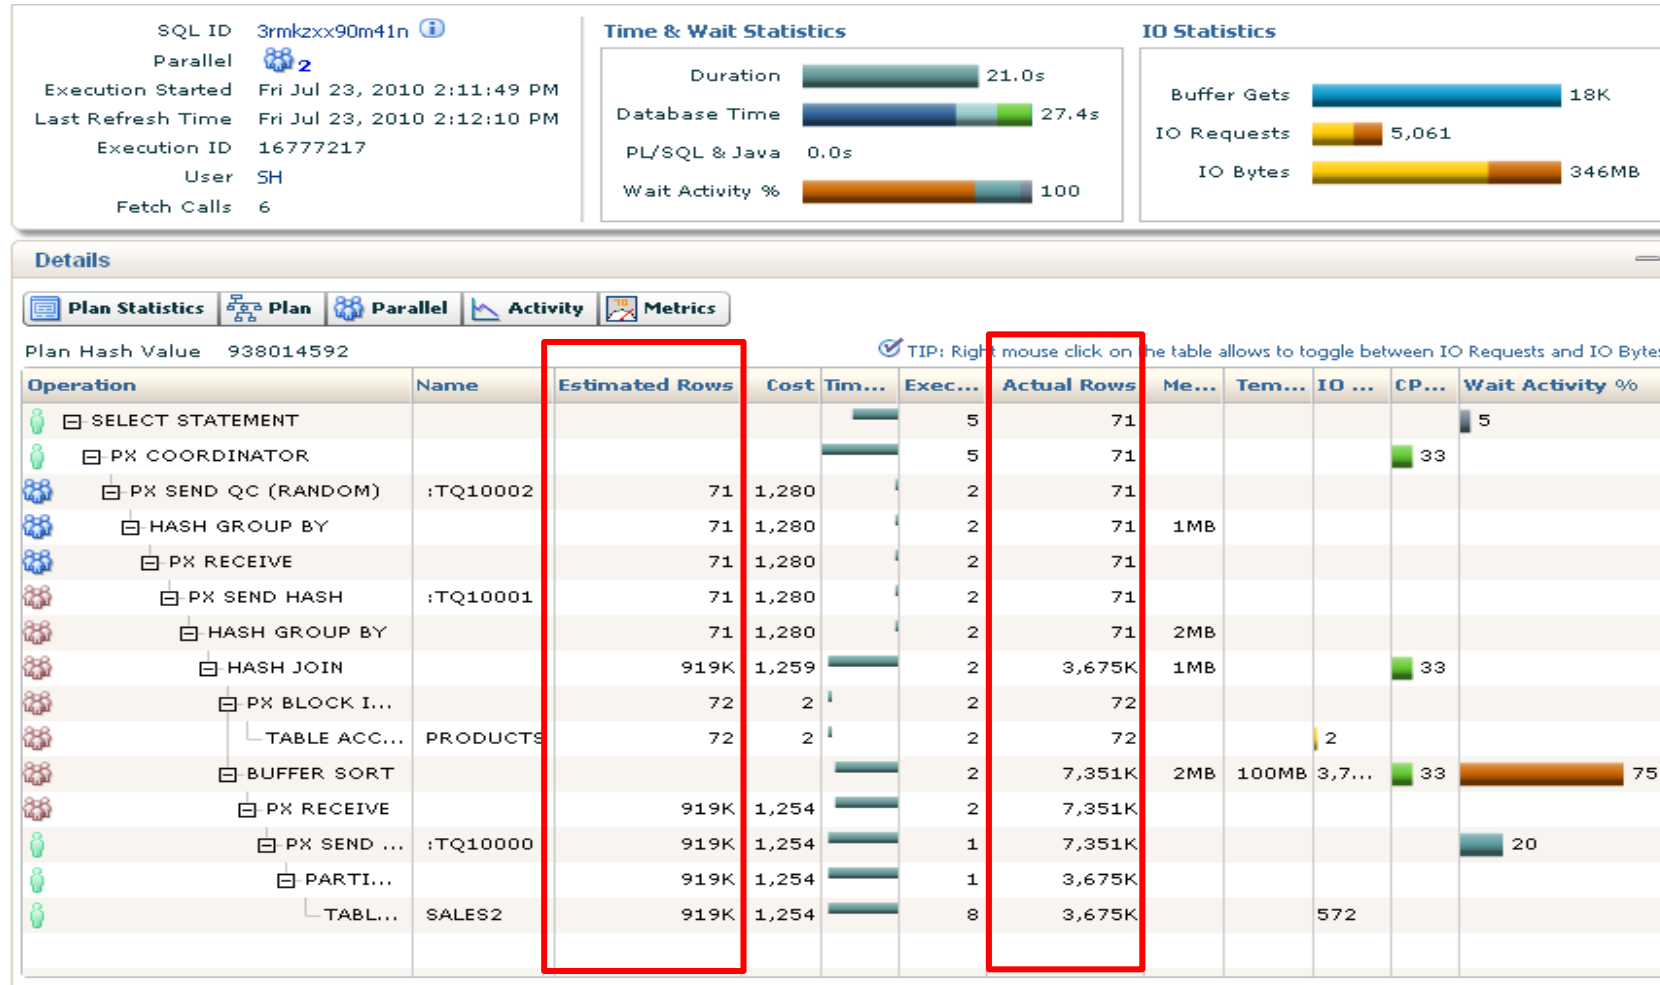
## Optimizer / Execution Plan

- Cardinality on the object and on the join level is determined by the optimizer to find the best execution plan
- For the optimizer that means number of rows returned by an operation
- The column **ROWS** in the execution plan or **ESTIMATED ROWS** in sql monitor shows this information
- Correct information is crucial for correct execution plans
- Cardinality feedback (object level):
  - optimizer estimates cardinality
  - actual rows processed are kept in row source tree
  - If estimations are wrong by factor 2 or more they are overwritten






# SQL Statements / Cardinality using EM

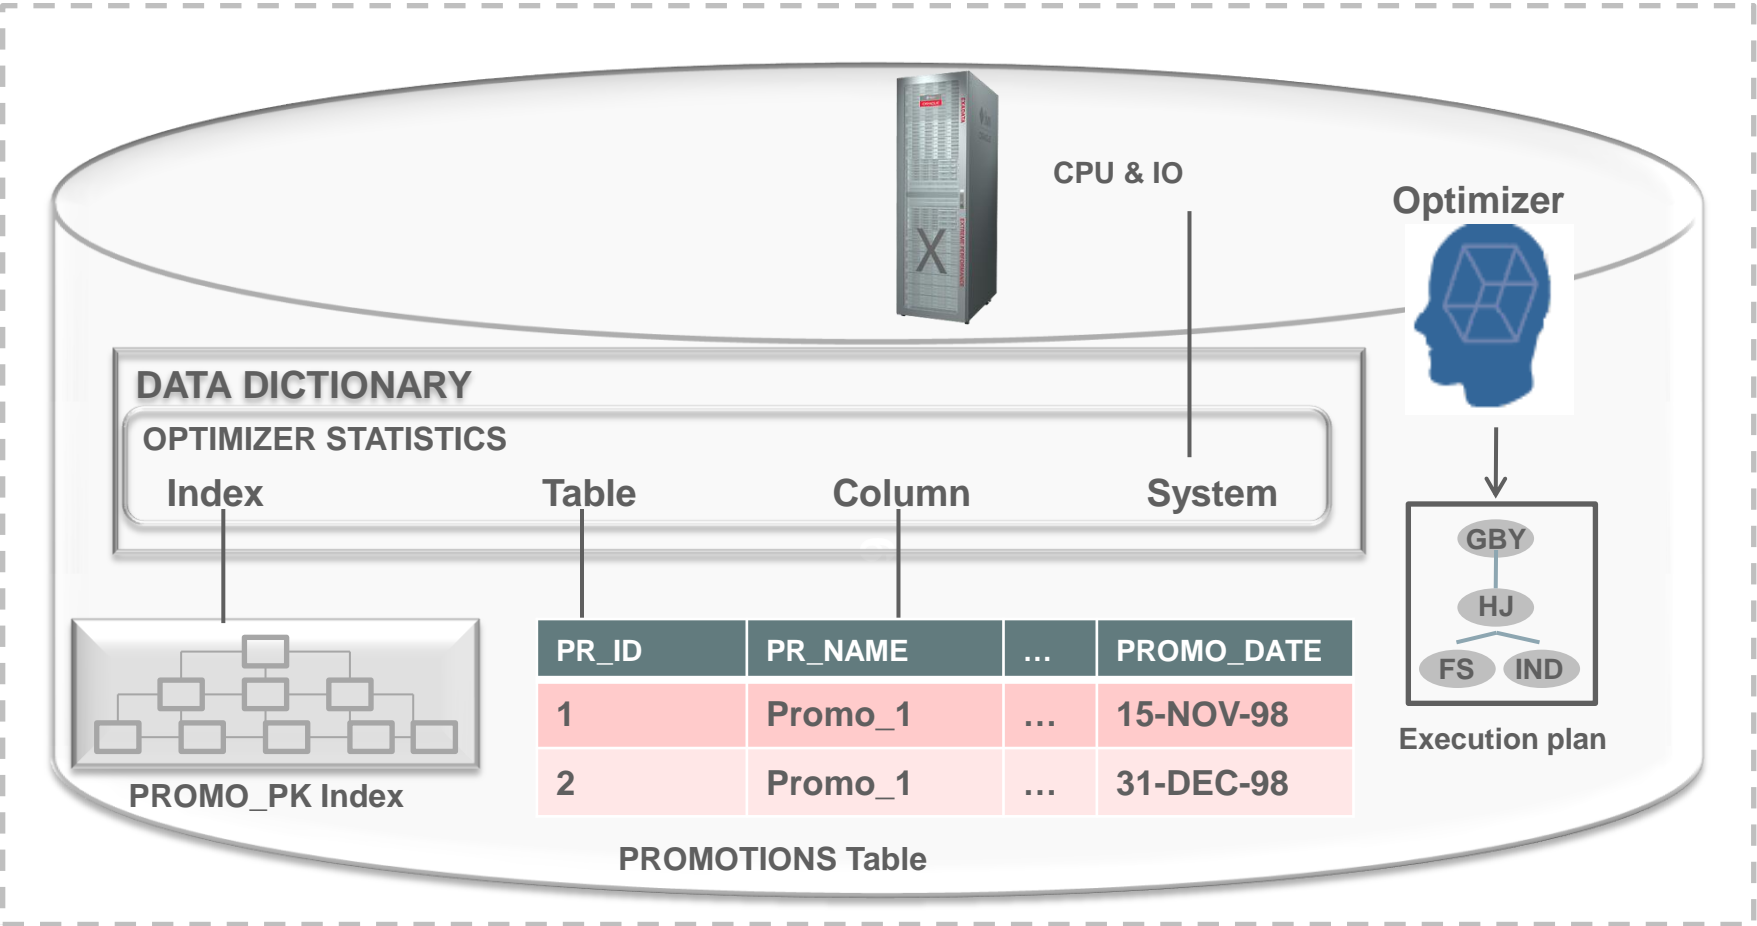


# SQL Statements / Incorrect Cardinality


## Optimizer / Execution Plan



# Optimizer Statistics







# Optimizer Statistics / Stale Statistics

- Statistics are considered stale when 10% or more of the rows in the object have changed
  - Changes include, inserts, updates, deletes etc.
- Query dictionary to check if statistics are stale

```
SELECT table_name, stale_stats FROM user_tab_statistics;
```

Table Name	Stale Statistics
Sales	NO
Customers	Yes
Product	-



No means stats are good




Yes means stats are stale



Null means no stats

Solution gather statistics



# Optimizer Statistics / No Statistics

PLAN\_TABLE\_OUTPUT

Plan hash value: 312224462

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		3844	326K	1252 (1)	00:00:16		
1	PARTITION HASH ALL		3844	326K	1252 (1)	00:00:16	1	8
* 2	TABLE ACCESS FULL	SALES2	3844	326K	1252 (1)	00:00:16	1	8


Predicate Information (identified by operation id):

2 - filter("CUST\_ID">=5 AND "CUST\_ID"<=30)

Note


- dynamic sampling used for this statement (level=2)

Solution gather statistics



# Optimizer Statistics / No Statistics – Dynamic Sampling

- Optimizer gathers statistics during parse operation
  - *alter system set optimizer\_dynamic\_sampling = 6;*
- Optimizer Dynamic Sampling (no joins):
  - 0 = off
  - 2 = *DEFAULT, tables without stats, 32 blocks*
  - 3 = *level 2 + complex single predicates, 64 blocks*
  - 4 = *even if table is analyzed, statistics are gathered*
  - 5 = *level 4, 2x default blocks*
  - 6 = *level 4, 4 x default blocks → PQ*
  - 7 = *level 4, 8 x default blocks*
  - 8 = *level 4, 32 x default blocks*
  - 9 = *level 4, 128 x default blocks*
  - 10 = *level 4, complete table*




# Optimizer Statistics / How to gather Statistics

- Analyze command is deprecated
- The GATHER\_\*\_STATS procedure takes 13 parameters
  - *Schema Name*
  - *Table Name*
  - *Partition Name*
  - ...
- Gather Statistics command should be simple
- From 11g onwards use default estimate\_percent  
*AUTO\_SAMPLE\_SIZE*

```
SQL> BEGIN
  2  dbms_stats.gather_table_stats('SH','SALES');
  3  END;
  4  /
```


PL/SQL procedure successfully completed.



# Optimizer Statistics / *AUTO\_SAMPLE\_SIZE*

Run Num	AUTO_SAMPLE_SIZE	10% SAMPLE	100% SAMPLE
1	00:02:21.86	00:02:31.56	00:08:24.10
2	00:02:38.11	00:02:49.49	00:07:38.25
3	00:02:39.31	00:02:38.55	00:07:37.83

Column Name	NDV with AUTO_SAMPLE_SIZE	NDV with 10% SAMPLE	NDV with 100% SAMPLE
C1	59852	31464	60351
C2	1270912	608544	1289760
C3	768384	359424	777942



# Optimizer Statistics / Example of Data Skew

```
SELECT * FROM HR.Employee WHERE Job_id = AD_VP;
```

NAME	ENUM	JOB
Kochhar	101	AD_VP
De Haan	102	AD_VP




HR Employee table

Last name	Em id	Job id
SMITH	99	CLERK
ALLEN	7499	CLERK
WARD	2021	CLERK
KOCHHAR	101	AD_VP
De Haan	102	AD_VP
CLARK	7782	CLERK

Optimizer assumes even distribution

Cardinality estimate is  $\frac{\text{NUM\_ROWS}}{\text{NDV}} = \frac{107}{19} = 6$

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time
0	SELECT STATEMENT		1	2	2	00:00:00.01
* 1	TABLE ACCESS FULL	EMPLOYEES	1	6	2	00:00:00.01




# Optimizer Statistics / Solution: Histograms

```
EXEC DBMS_STATS.GATHER_TABLE_STATS('HR','EMPLOYEES',method_opt =>
'FOR ALL COLUMNS SIZE SKEWONLY');
```

```
SELECT column_name
       , num_distinct
       , histogram
FROM user_tab_col_statistics
WHERE table_name =
      'EMPLOYEES';
```


COLUMN_NAME	NUM_DISTINCT	HISTOGRAM
EMPLOYEE_ID	107	NONE
FIRST_NAME	91	NONE
LAST_NAME	102	NONE
EMAIL	107	NONE
PHONE_NUMBER	107	NONE
HIRE_DATE	98	NONE
JOB_ID	19	FREQUENCY
SALARY	58	NONE
COMMISSION_PCT	7	NONE
MANAGER_ID	18	FREQUENCY
DEPARTMENT_ID	11	FREQUENCY



# Optimizer Statistics / Multiple correlated Columns

SELECT ... WHERE model = '530xi' and make = 'BMW';

Make	Model	Color
BMW	530xi	red
BMW	530xi	black
BMW	530xi	silver




Vehicles Table

Make	Model	Color
BMW	530xi	RED
BMW	530xi	BLACK
BMW	530xi	SILVER
PORSCHE	911	RED
MERC	SLK	RED
MERC	C320	SLIVER

Cardinality est. is    ROWS (12)     $\frac{1}{\text{NDV C1 (3)}} \times \frac{1}{\text{NDV C2 (4)}} = 1$

Id	Operation	Name	Starts	E-Rows	A-Rows
0	SELECT STATEMENT		1		1
1	SORT AGGREGATE		1	1	1
* 2	TABLE ACCESS FULL	VEHICLES	1	1	3





# Optimizer Statistics / Solution: Extended Statistics


*Create extended statistics on the Model & Make columns*

*exec dbms\_stats.gather\_table\_stats('DWH', 'VEHICLES', degree=>8,  
method\_opt=> 'for columns (MODEL, MAKE) size 256 ');*

```
SELECT column_name  
       , num_distinct  
       , histogram  
FROM user_tab_col_statistics  
WHERE table_name = 'VEHICLES';
```


COLUMN_NAME	NUM_DISTINCT
-----	
MAKE	5
MODEL	9
COLOR	7
SYS_STU80PK2S\$PEWHARK2CP3#1F#G	9


New Column with system  
generated name



# Optimizer Statistics / Solution: Adaptive Query Plan

- The optimizer in 12g checks the cardinality during execution
- Results are buffered before processed





# Optimizer Statistics / Solutions

Cause	Solution
Stale or missing statistics	DBMS_STATS
Data Skew	Create a histogram
Multiple single column predicates on a table	Create a column group using DBMS_STATS.CREATE_EXTENDED_STATS
Multiple columns used in a join	Create a column group using DBMS_STATS.CREATE_EXTENDED_STATS
Function wrapped column	Create statistics on the function wrapped column using DBMS_STATS.CREATE_EXTENDED_STATS
Complicated expression containing columns from multiple tables	Use dynamic sampling level 4 or higher

Active in last 2 hours							Refresh Manual Refresh		
Status	Duration	SQL ID	User	Parallel	Database Time	IO Requests	Start	Ended	SQL Text
	6.0s	8wbxw9vvq98a7	DWH_DATA	64  8	57.2s	8513	9:27:10 AM		SELECT /*# QUERYLOG #*/ /*+ ORDERED USE_H
	1.2m	9fcpzwgsskxzd	DWH_DATA	64  8	25.9m	186K	9:21:17 AM	9:22:28 AM	SELECT /*+ ORDERED USE_HASH(t1) INDEX(t1) PUS
	7.0s	572fbaj0fdw2b	SYS		6.4s	222	9:17:25 AM	9:17:32 AM	select output from table(dbms_workload_repository.awr
	7.0s	572fbaj0fdw2b	SYS		6.8s	220	9:17:17 AM	9:17:24 AM	select output from table(dbms_workload_repository.awr
	7.0s	572fbaj0fdw2b	SYS		6.9s	232	9:17:09 AM	9:17:16 AM	select output from table(dbms_workload_repository.awr
	7.0s	572fbaj0fdw2b	SYS		7.2s	224	9:17:00 AM	9:17:07 AM	select output from table(dbms_workload_repository.awr
	7.0s	572fbaj0fdw2b	SYS		6.4s	213	9:16:52 AM	9:16:59 AM	select output from table(dbms_workload_repository.awr
	7.0s	572fbaj0fdw2b	SYS		6.6s	215	9:16:44 AM	9:16:51 AM	select output from table(dbms_workload_repository.awr
	7.0s	572fbaj0fdw2b	SYS		6.8s	265	9:16:36 AM	9:16:43 AM	select output from table(dbms_workload_repository.awr
	12.0s	572fbaj0fdw2b	SYS		11.0s	970	9:16:23 AM	9:16:35 AM	select output from table(dbms_workload_repository.awr
	7.0m	9fcpzwgsskxzd	DWH_DATA	64  8	2.2h	2374K	9:01:42 AM	9:08:39 AM	SELECT /*+ ORDERED USE_HASH(t1) INDEX(t1) PUS
	1.2m	6xr88j72tap5v	DWH_DATA	64  8	41.9m	73K	9:00:25 AM	9:01:37 AM	SELECT /*+ ORDERED USE_HASH(t1) INDEX(t1) PUS

**Overview**

SQL ID: 8wbxw9vvq98a7

Parallel: 64

Execution Started: Wed Apr 7, 2010 9:27:10 AM

Last Refresh Time: Wed Apr 7, 2010 9:29:04 AM

Execution ID: 16777216

User: DWH\_DATA

Fetch Calls: 0

**Time & Wait Statistics**

Duration: 2.0m

Database Time: 29.6m

Wait Activity %: 100

PL/SQL & Java: 0.0s

**IO Statistics**

IO Requests: 282K

IO Bytes: 34G

Buffer Gets: 137K

Cell Offload Efficiency: -96.08%

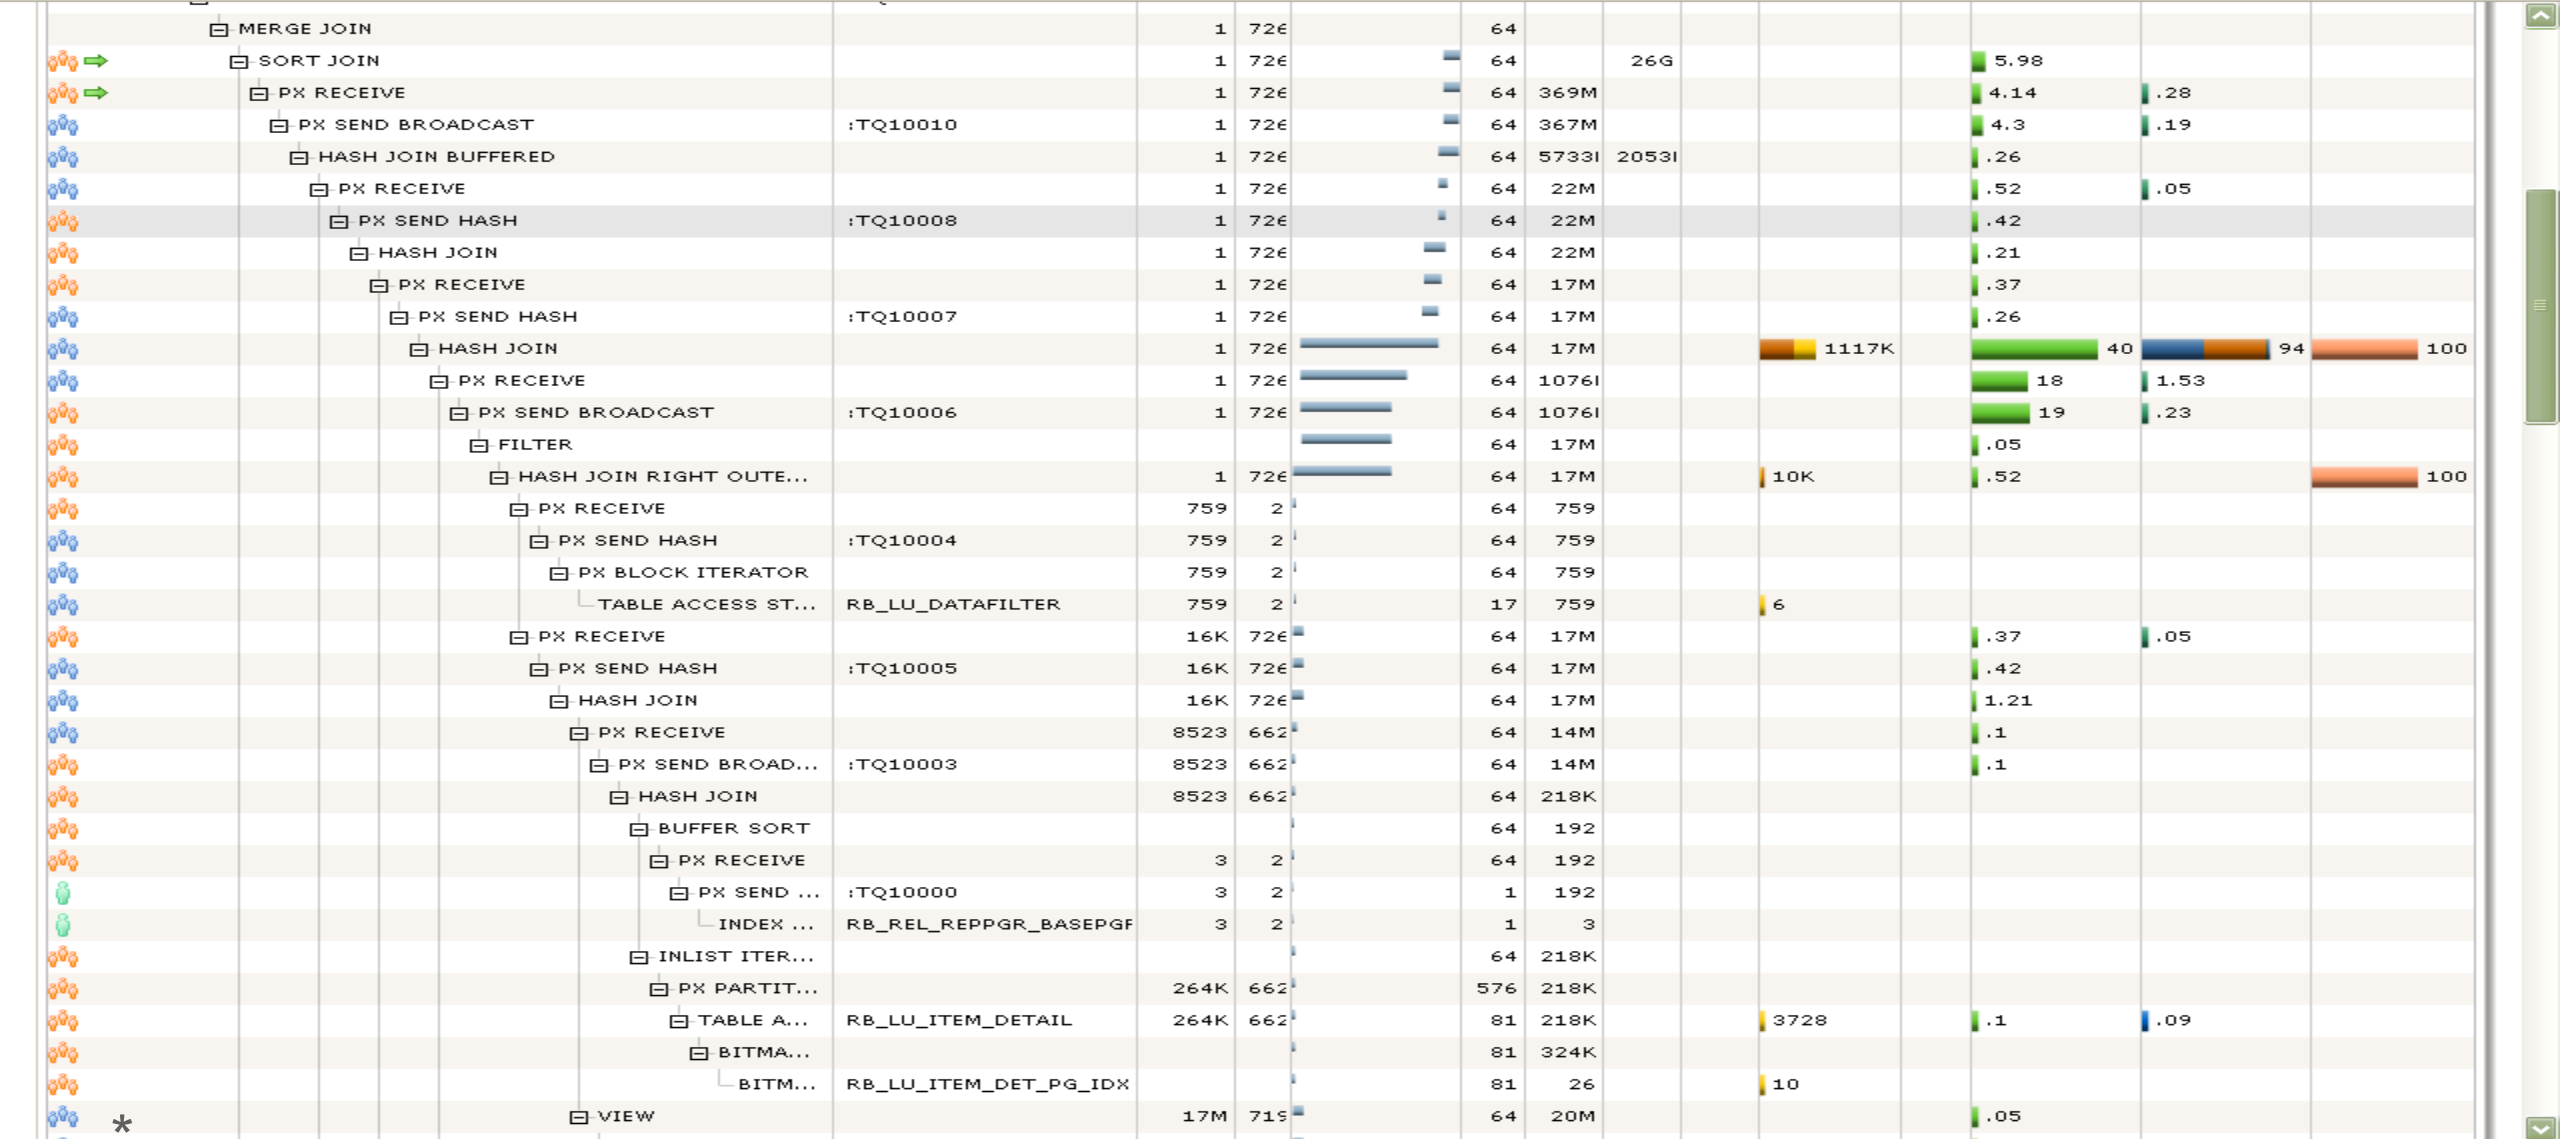
**Details**

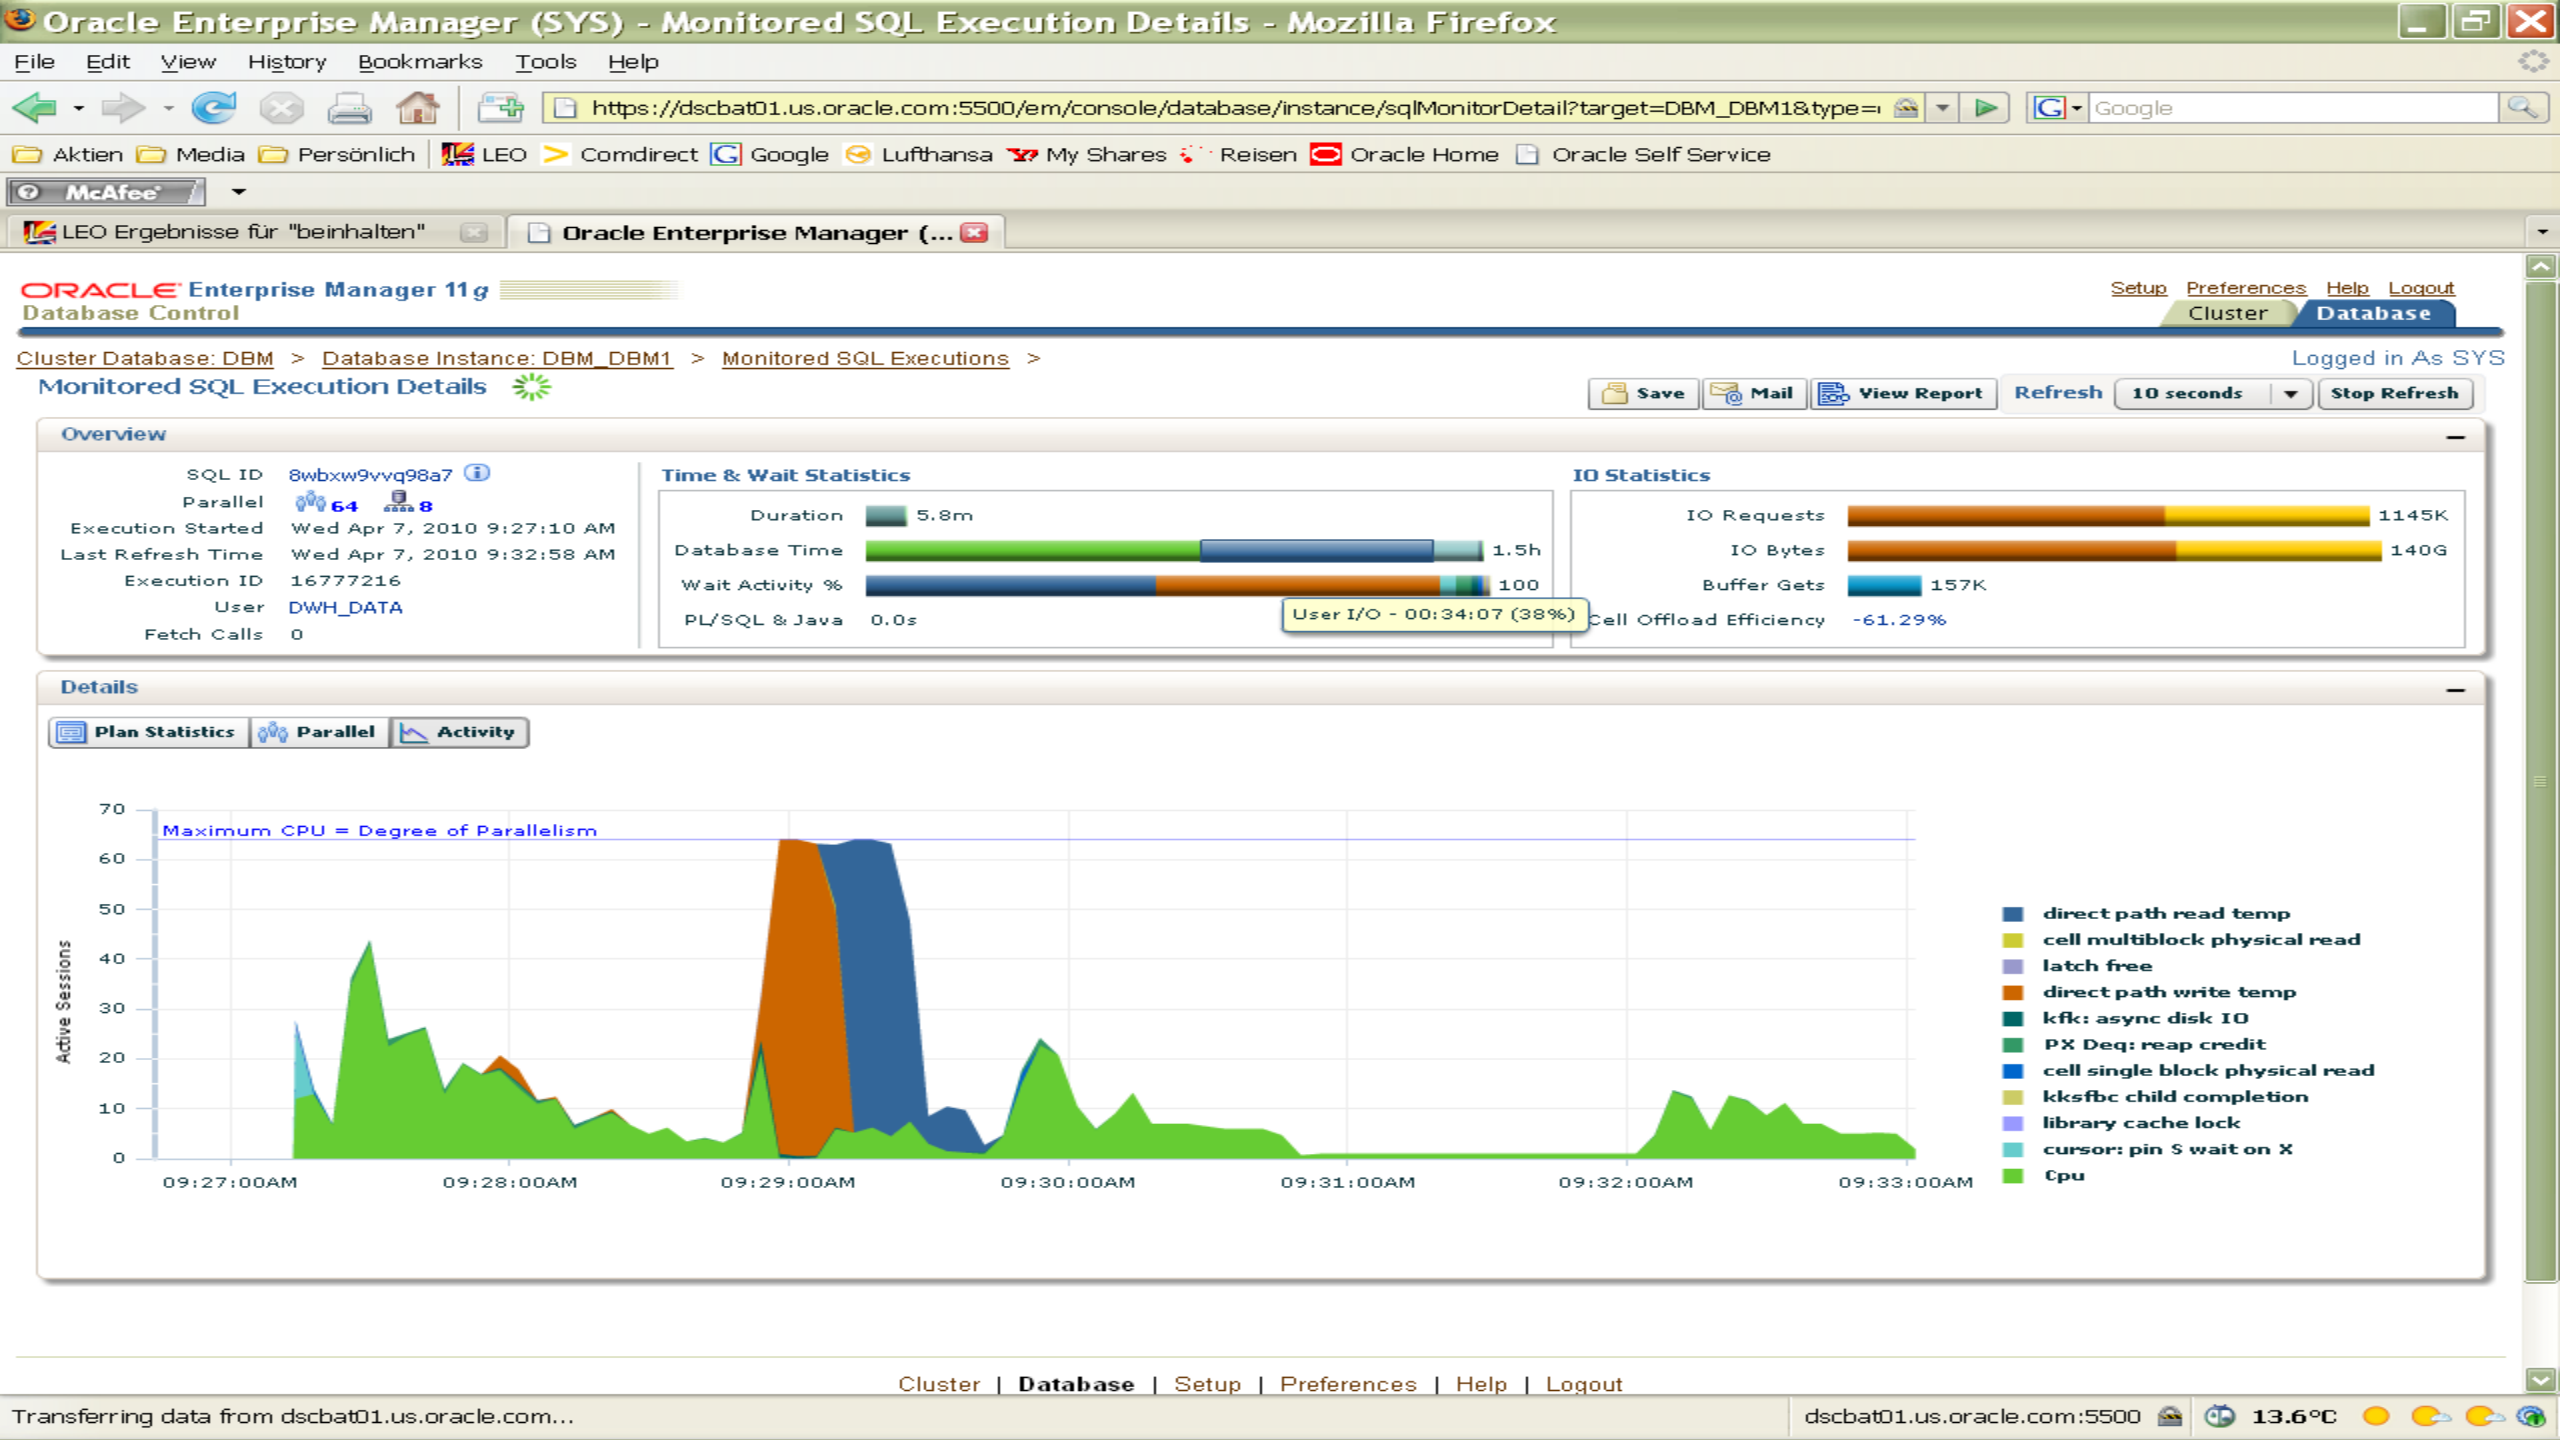
Plan Statistics Parallel Activity

Plan Hash Value: 233075956

TIP: Right mouse click on the table allows to toggle between IO Requests and IO Bytes

Operation	Name	Estimat...	Cost	Timeline(117s)	Exe...	Actu...	Mem...	Temp	IO Requests	Cell ...	CPU Activity %	Wait Activity %	Progress %
SELECT STATEMENT					129								
PX COORDINATOR					129								
PX SEND QC (RANDOM)	:TQ10012	1	726K										
HASH GROUP BY		1	726K										
PX RECEIVE		1	726K										
PX SEND HASH	:TQ10011	1	726K										
MERGE JOIN		1	726K										
SORT JOIN		1	726K										
PX RECEIVE		1	726K										
PX SEND BROADCAST	:TQ10010	1	726K										
HASH JOIN BUFFERED		1	726K										
PX RECEIVE		1	726K										
PX SEND HASH	:TQ10008	1	726K										
HASH JOIN		1	726K										
PX RECEIVE		1	726K										
PX SEND HA...	:TQ10007	1	726K		64								
HASH JOIN		1	726K		64		55G	44G	310K		44	90	
PX RECE...		1	726K		64	1074M					24	3.08	

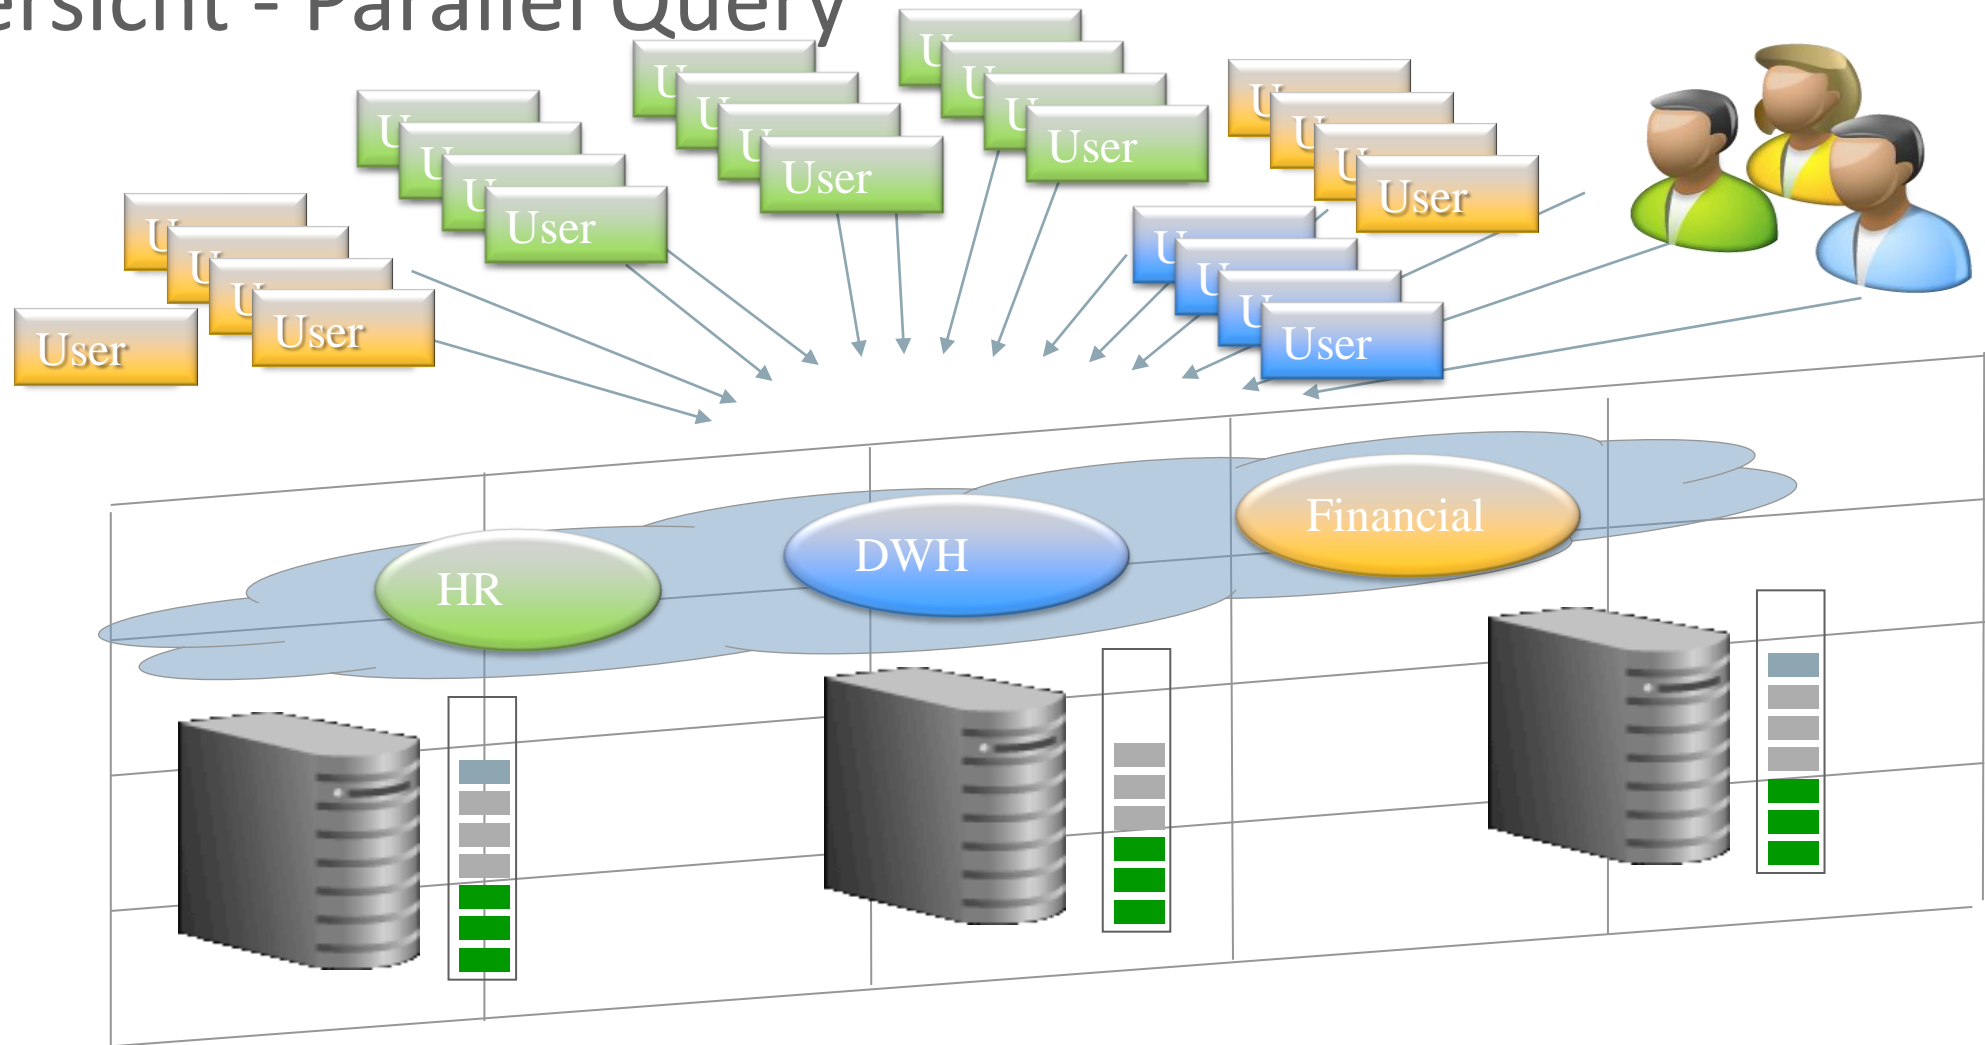








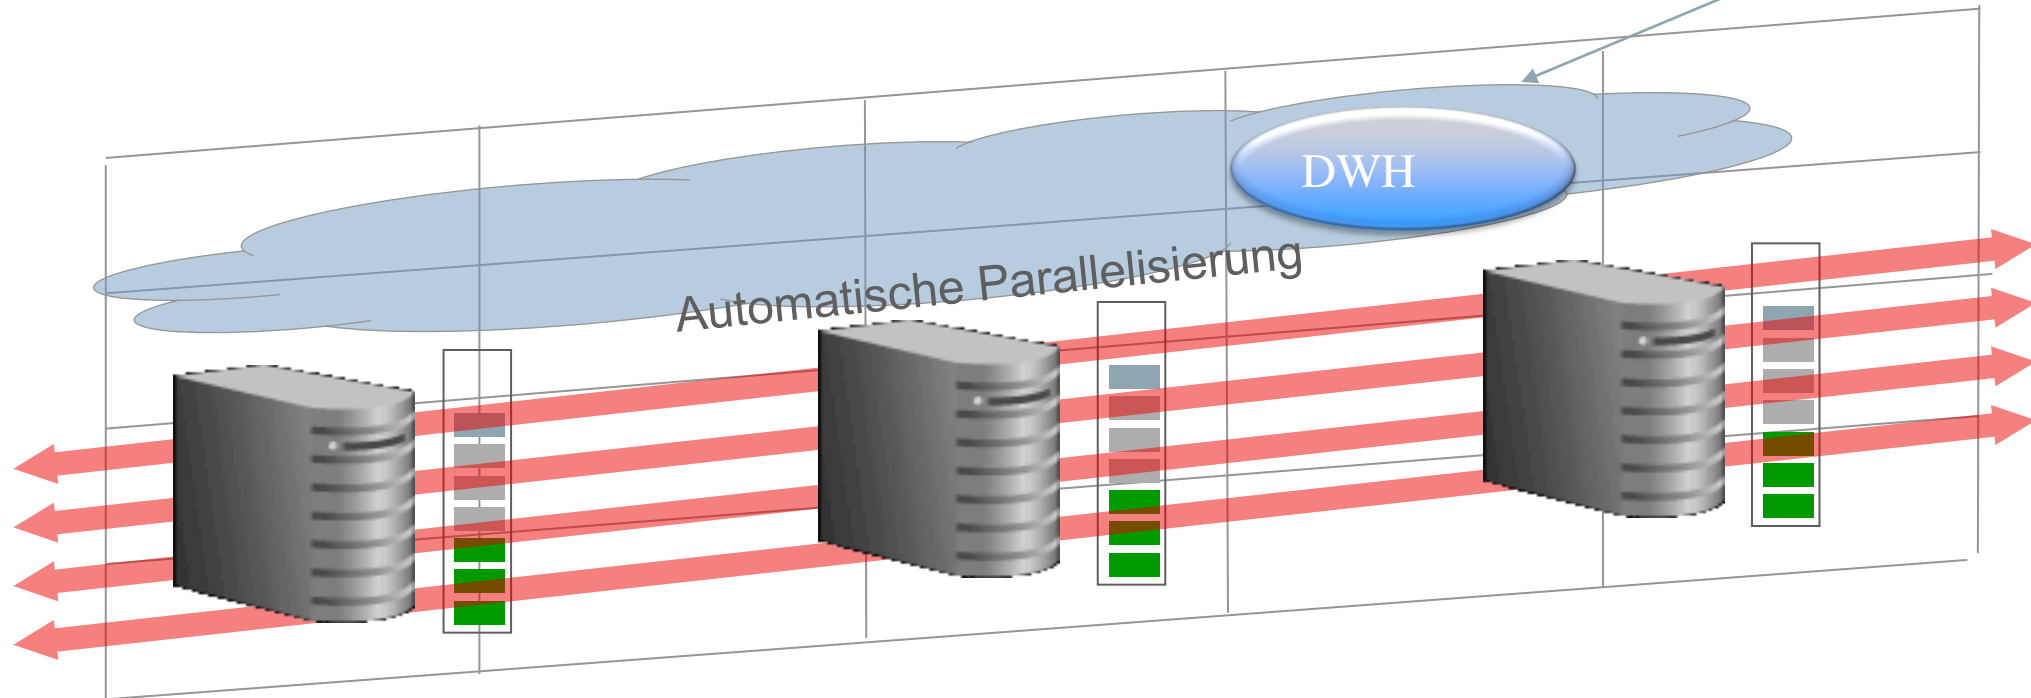
# Übersicht - Parallel Query




# Übersicht - Parallel Query



Oracle ist in der Lage, sämtliche Operationen innerhalb eines Knotens oder zwischen den Knoten zu parallelisieren





# PO: Producers & Consumers

SQL ID fgs5yggzfr38h ⓘ  
Parallel 8 8  
Execution Started Mon Jul 19, 2010 1:48:03 PM  
Last Refresh Time Mon Jul 19, 2010 1:50:17 PM  
Execution ID 33554432  
User RETAIL  
Fetch Calls 1

Details

Plan Statistics Plan Parallel Activity

Plan Hash Value 1108174806

Operation	Name
SELECT STATEMENT	
PX COORDINATOR	
PX SEND QC (ORDER)	:TQ1
SORT ORDER BY	
PX RECEIVE	
PX SEND RANGE	:TQ1
HASH JOIN BUFFERED	
PX RECEIVE	
PX SEND HASH	:TQ1
VIEW	
WINDOW SORT	
PX RECEIVE	
PX SEND RANGE	:TQ1
HASH GROUP BY	
PX RECEIVE	
PX SEND H...	:TQ1
HASH GR...	
HASH J...	
PX RE...	
PX ...	:TQ1
H...	
P...	
P	:TQ1

Transferring data from database to us.oracle.com...

Overview

SQL ID fgs5yggzfr38h ⓘ  
Parallel 8 8  
Execution Started Mon Jul 19, 2010 1:48:03 PM  
Last Refresh Time Mon Jul 19, 2010 1:49:59 PM  
Execution ID 33554432  
User RETAIL  
Fetch Calls 0

Tin  
D.  
F  
W

Details

Plan Statistics Plan Parallel Activity

☒ Show Instance Nodes

Parallel Server	Database Time
All Parallel Servers	
Instance 2	
Parallel Coordinator	0.3s
Parallel Set 1	43.4s
Parallel Set 2	37.6s
Instance 1	
Parallel Set 1	43.0s
Parallel Set 2	36.1s
Instance 8	1.4m
Instance 7	1.5m
Instance 6	1.4m
Instance 5	1.4m
Instance 4	1.4m
Instance 3	1.4m

Oracle Enterprise Manager (SYS) - Monitored SQL Execution Details - Mozilla Firefox

FileEditViewHistoryBookmarksToolsHelp

←→↺↻🖨🏠➕🔍

https://dscbat01.us.oracle.com:5500/em/console/database/instance/sqlMonitorDetail?target=DBM\_DBM1&type=

Google

AktienMediaPersönlichLEOComdirectGoogleLufthansaMy SharesReisenOracle HomeOracle Self Service

McAfee

LEO Ergebnisse für "beinhalten"Oracle Enterprise Manager (...)

ORACLE Enterprise Manager 11gDatabase Control

SetupPreferencesHelpLogoutClusterDatabase

Cluster Database: DBM > Database Instance: DBM\_DBM1 > Monitored SQL Executions >

Monitored SQL Execution Details

SaveMailView ReportRefresh10 secondsStop Refresh

Overview

SQL ID8wbxw9vvq98a7

Parallel648

Execution StartedWed Apr 7, 2010 9:27:10 AM

Last Refresh TimeWed Apr 7, 2010 9:31:45 AM

Execution ID16777216

UserDWH\_DATA

Fetch Calls0

Time & Wait Statistics

Duration4.6m

Database Time1.3h

Wait Activity %100

PL/SQL & Java0.0s

IO Statistics

IO Requests1145K

IO Bytes140G

Buffer Gets157K

Cell Offload Efficiency-61.29%

Details

Plan StatisticsParallelActivity

☒ Show Instance Nodes

Parallel Server

Database Time

Wait Activity %

IO Requests

Cell Offload ...

Buffer Gets

All Parallel Servers

Instance 1

Parallel Coordinator0.4s

Parallel Set 1

Parallel Server 9 (p000)59.2s1.3917K-63.93309

Parallel Server 10 (p001)1.0m1.4917K-63.938529

Parallel Server 11 (p002)1.2m1.4917K-63.93393

Parallel Server 12 (p003)1.1m1.4417K-63.937912

Parallel Server 13 (p004)57.7s1.3917K-63.932544

Parallel Server 14 (p005)56.0s1.3517K-63.93313

Parallel Server 15 (p006)57.0s1.2517K-63.936701

Parallel Server 16 (p007)59.4s1.3517K-63.936593

Parallel Set 21.4m.09368

Instance 2

Parallel Set 18.2m13149K

Parallel Set 252.7s.463693

Instance 89.1m12135K34K

Instance 79.5m12137K20K

Waiting for dscbat01.us.oracle.com...dscbat01.us.oracle.com:550013.6°C



# Performance Zusammenfassung / Empfehlungen

## Analyse

- Analyse des Gesamtsystems notwendig, bevor Maßnahmen ergriffen werden
- Konzentrieren der Maßnahmen auf die Zeitfresser

## Applikation

- Was für ein Typ ist die Applikation? OLTP, DWH, Mix
- Nutzt die Applikation die Ressourcen der Datenbank und des Systems?
- Monitoring der Applikation, nicht nur in der Datenbank

## SQL Statements

- Möglichkeiten zur Analyse/Monitoring von SQL Statements
- Erstellen von aktuellen Statistiken
- Tuning von SQL Statements / optimaler Execution Plan / Parallel Query

## Systemperformance

- Gute Zusammenarbeit aller beteiligten Abteilungen
- Well Balanced System